

02-613 Week 9

Algorithms and Advanced Data Structures

Aidan Jan

October 29, 2025

Suffix Tries

Problem: Given a large known, fixed **text**, e.g. a dictionary, and many unknowns of changing search queries, how do we design a data structure that can quickly search up terms in the structure, but also substrings?

- For example, if “uninterested” was in the dictionary, and we wanted to search for “interested” or “interest”, they should both show up.
- Using any of our structures so far, there is not an efficient way to do this.
- What we *can* do is make a tree, where each node represents a letter, where every suffix is a path from the root to some leaf node.
- We will use the string $S = ABAABA\$$. ($\$$ is a terminus character to mark the end of the string.)

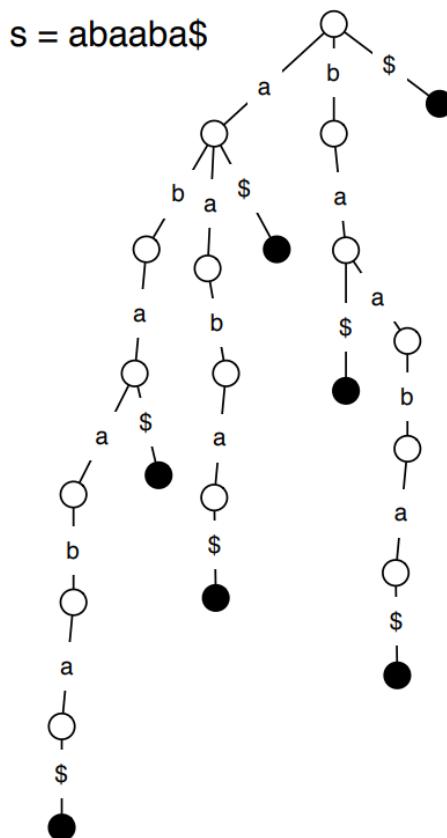


Image by Dr. Carl Kingsford @CMU

This structure allows for searching for any word in our dictionary, but also any suffix.

- This is known as a **trie**. (Not a tree, since trees don't allow internal nodes with only one child.)

Runtime of Tries

The runtime of a trie is $\Theta(\Sigma)$, where Σ is the size of our alphabet. For the English alphabet, $\Sigma = 27$ (26 letters, plus an end-of-word marker, $\$$.)

- Similarly, a DNA sequence would have $\Sigma = 5$ since it includes $A, C, T, G, \$$.
 - We add a $\$$ to be able to distinguish when a valid string ends. If we want to find if a query is a suffix of our text, then the terminators become very important.

Searching in Tries

To search for a term q in our text T , we simply follow the path down from the root, since every substring would have a unique path from the root.

- If we want to know if q is a suffix of T , we can do that too since q would be a substring of T .
 - If we want to know the number of times q is in T , then we follow the path down for q , and at the last character, count the number of leaf nodes in the subtree.
 - To find the longest repeated substring, we return the deepest branching node.
 - To find the lexicographically smallest suffix, from the root we follow the tree down always following the branch with the smallest character, assuming $\$$ is less than any other character in our alphabet.

Suffix Trees

Consider the string $s = a^n b^n$, where $n \in \mathbb{Z}$. Basically, a library of string containing a series of a followed by a suffix if b . The trie would look like:

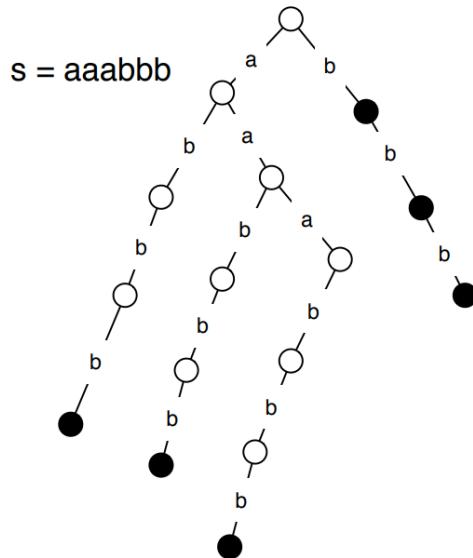


Image by Dr. Carl Kingsford @CMU

In this case, once we get into the string of b 's, there is only one long branch. However, if we wanted to encode something like the human genome (with 3B base pairs), we can't store $O(n^2)$ nodes like we do here. To make this more space efficient, we can compress this.

Key Points of Construction

- The number of leaves we have after compression is $O(n)$ with the above example.
- All internal nodes are branching.
- If instead we store string indices rather than the actual string on every edge, we don't have to build every substring in our tree either, which brings storage down to $O(n)$.
- Tries are useless in practice, we usually build the suffix tree directly.

Building a Suffix Tree

- Start from the empty string (root)
- Add one character at the beginning of the string, and add the suffix link too
- Add the next character to the end of every suffix following the suffix link back to the root
 - This is equivalent to adding the next character to the end of every existing suffix.
- Repeat until entire string is represented in the tree.

```
current suffix = longest suffix
do:
    add child labeled S[i] to current suffix
    follow suffix link to set current suffix to next
    add suffix link to new node
while not every letter of string is in the tree
```

For this class, construction is not discussed much. Use Ukkonen's Algorithm for construction in $O(n)$ time.

Finding Substrings

Suppose we want to find the longest common substring between S and a query.

- First, let's find the longest common prefix. To do this, we simply start with q , and walk down the suffix tree as far as we can.
- Now, extending this to all substrings, we can change the "starting point" by following the suffix links.
 - If we get to a point where we cannot extend our prefix, we follow the last suffix link (which effectively moves the starting point forward by a few positions), and continue iterating!
 - This way, the algorithm will take at longest $O(m)$, where m is the length of the query, since we never backtrack in the query string. This beats the naive method which is $O(m^2)$.

For example, suppose we want to find the longest common substring between $q = BBAA$ and $S = ABAABA$

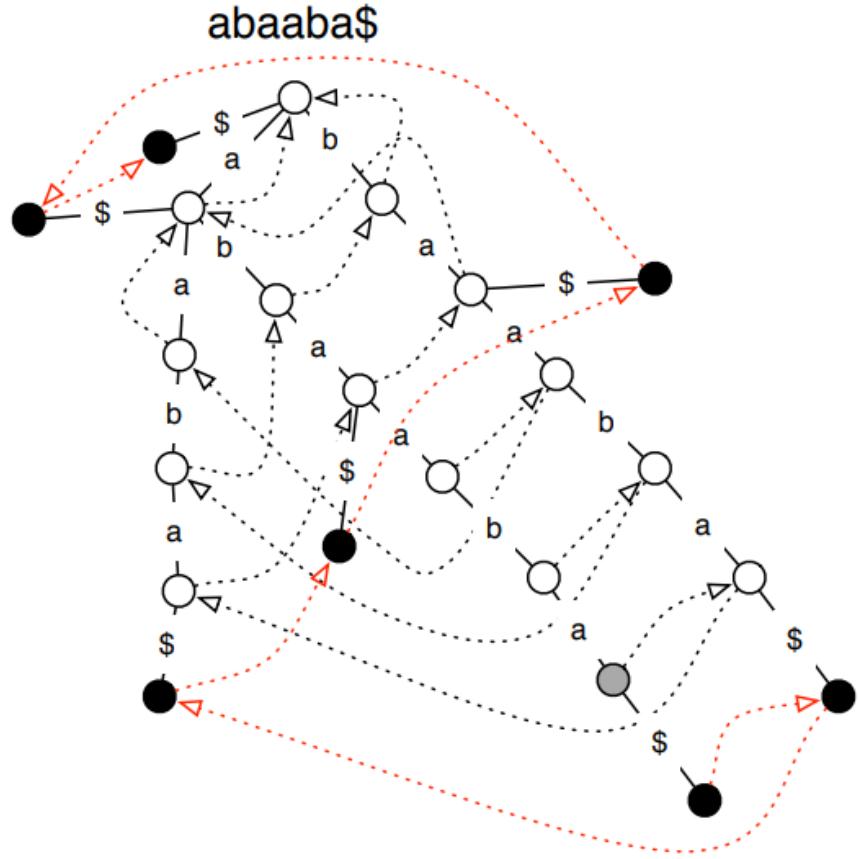


Image by Dr. Carl Kingsford @CMU

Suffix Tree of Multiple Strings

Given a set $S = \{s_1, s_2, \dots, s_n\}$ sequences from an alphabet, represent as a suffix tree. How do we construct this?

- There is an easy way: concatenate all the strings.
- Consider $S = \{aat, tag, gat\}$. We can build an $S' = aat\$_1tag\$_2gat\$_3$.
- Now, build a suffix tree using Ukkonen's on S' , iterate through each node, and delete everything past the first terminal character.
- This runs in $O(n)$, where n is the total length of all the sequences.

For the above example, we will get the result:

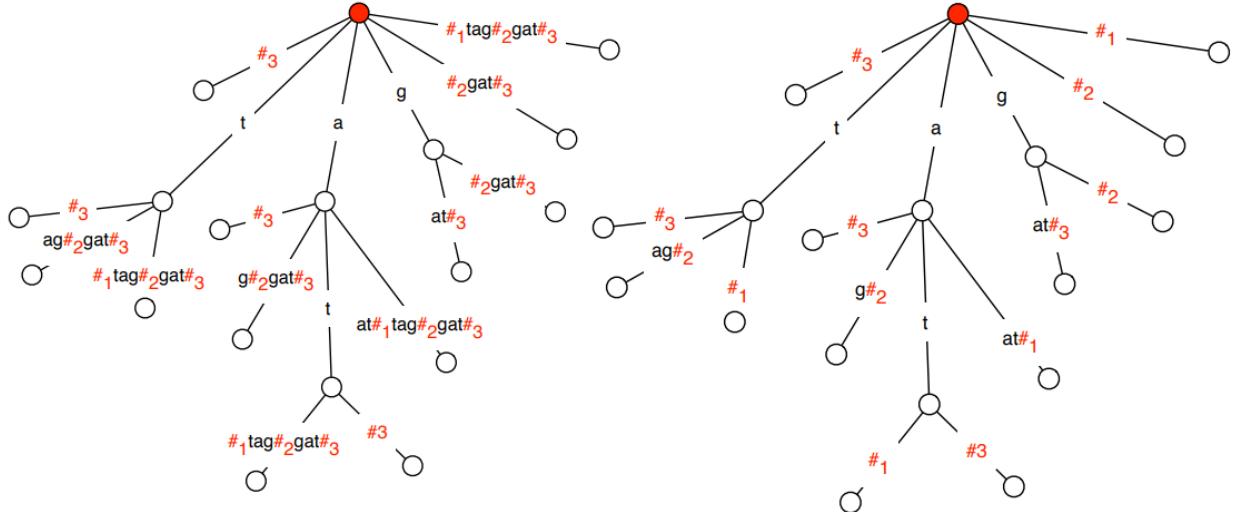


Image by Dr. Carl Kingsford @CMU

Longest Common Substring

Building upon the suffix tree of multiple strings, how do we find the longest common substring between two sequences?

- We can do this also in $O(n)$ time, where n is the length of the two strings combined.
- Build a suffix tree containing both strings of interest, and run a DFS.
- Propagate the value of each leaf node (always a terminus, $\$_i$) backwards, and the deepest node that can reach both terminus would be the longest common substring.

We can also use this strategy to find all the strings in a database $\{S_1, S_2, \dots, S_n\}$ that contain any query, q .

Longest Common Extension

Suppose we are given strings S and T . We will be provided with many pairs (i, j) as queries, and we want to be able to quickly find the longest substring of S starting at i that matches a substring of T starting at j .

- Build a suffix tree containing both S and T .
- Preprocess the tree so that the lowest common ancestors (LCA) can be found in constant time. (We won't prove this, but it can be done in $O(N)$, where N is the total length of the two strings.)
- Create an array mapping suffix numbers to leaf nodes.
- Given query (i, j) ,
 - Find the leaf nodes for i and j
 - Return string of the LCA for i and j
- Query can be found in $O(1)$.

LCE for Palindromes

A palindrome has the property where the left half is equal to the reverse of the right half. To find **maximal** palindromes in S (e.g., given a center point, find the longest palindrome that contains it), we construct the reverse string, S^r .

- Notice that if given a center point i , the right half would be a substring beginning at i on the forward string, and the left half would be a string starting at $(n - i)$ on the reverse string.
- With this, we can create a suffix tree containing S and S^r , find the LCE given point i , and the string going from the root to that intersection would be a palindrome.
- The total runtime is $O(n)$.

What if we allow mismatches in our palindrome?

k -mismatch Palindromes

This is the same problem, except this time, we allow for k -mismatches. For example, *TACOBAT* has 1 mismatch, since when comparing the left half to the reverse of the right half, there is one mismatch. To do this,

- Find the maximal palindrome from the center i . Suppose it has length $2j$, so the right of the string is at position $i + j$ on S , and the right begins at $n - i - j$ on S^r .
- For each k mismatches, we can add one "error" letter, and repeat the process with another LCE, with our query being $i + j + 1$ on S and $n - i - j - 1$ on S^r .
- Therefore, for k mismatches, it will take $O(kn)$.